




Figure 3: Our lying evaluation pipeline uses a three-step process anchored by a *proposition*, a variable capturing what we expect models to lie about and for which ground truth is known. (1) Prompt the model with two contrasting prompts – a pressure prompt designed to potentially induce deception and a neutral belief elicitation prompt. (2) Use LLMs to map both model statement (S) and belief (B) to their respective proposition values, with additional checks to ensure models consistently hold the belief B . (3) Measure lying by comparing statement S against belief B , as well as accuracy by comparing belief B against ground truth T . Example shortened for brevity.

Formally, let B denote the model’s belief (e.g., “the Great Wall was built primarily for symbolic purposes”) and let S represent the model’s response. We define lying as follows:

Lying is making a statement S where $S \neq B$, with the intent of causing a receiver to believe S is true.

Optionally, some definitions of lying also involve intent for the statement to mislead the recipient of the statement.

Evaluating beliefs. There is mounting evidence that AI models hold beliefs in a meaningful sense which affect their answers in a broad variety of ways, which affect their responses across different contexts (Meng et al., 2023; Nanda et al., 2023; Zou et al., 2023). For example, if a model believes that the capital of France is Paris, it will not only recall that fact when directly asked but also utilize it in related scenarios. For instance, if asked for directions from the south of France to the capital of France, the model would give directions to Paris as the destination. Thus, beliefs in this context refer to a model’s consistent internalization of facts, which guides the generation of its answers.

While the debate persists regarding whether these models truly “believe” things in the human sense, we operationalize this concept by focusing on consistency in response. We define a model’s belief as its answer to factual questions when not under any direct pressure to lie or be deceptive, and such that its answer is consistent with responses to related questions. Concentrating on the model’s underlying beliefs forms the foundation of our lying evaluation. See Section A.3 for additional discussion on belief consistency and how this opens the door to measuring lying.

Distinguishing honesty from accuracy. While honesty pertains to the intentionality behind the model’s output, accuracy is a measure of factual correctness. In most evaluations, a model’s factual accuracy tests against an objective ground truth label.

Inaccuracy is believing B where B is false.

Notably, a model could be highly accurate in its knowledge (i.e., aligning its beliefs with ground truth) but still be dishonest if it knowingly outputs false information. Conversely, a model could be less accurate but still be honest in its responses, as it does not intentionally mislead.

4 Model Alignment between Statements and Knowledge (MASK)

4.1 Dataset

MASK consists of 1,000 high-quality adversarially-selected human-labeled examples that, combined with our novel lying evaluation framework (Section 4.2), allow measurement of whether models are lying when exposed to situations that may induce deception. We also maintain a separate, held-out set of 500 examples to track potential overfitting. The results presented in this paper are based on the full set of 1,500 examples.

Each example in MASK consists of four components: a proposition, ground truth, pressure prompt, and belief elicitation prompt. The proposition is a simple statement with a pivotal variable that is binary or numerical (for example: “The US poverty rate in 2020 {was/was not} above 5%” or “The US poverty rate in 2020 was {value}”). The ground truth is a factually correct resolution of this proposition (e.g. “The US poverty rate in 2020 was 11.4%”). The pressure prompt is a human-crafted prompt designed to incentivize the model to make a false statement with respect to the proposition; these prompts present situations where honesty could conflict with other objectives, creating an incentive for the model to lie. The belief elicitation prompt is used to reveal the model’s actual belief about the proposition by asking about it in a neutral manner. Figure 5 includes three examples of each of these components.

To enable comprehensive evaluation of model lying, we collect pressure prompts across diverse scenarios that might trigger different patterns of deceptive behavior. In the Appendix, Table 2 defines the six distinct *archetypes* our prompts align with.


Model	$P(\text{Lie})$ ($\downarrow$)
Grok 2	63.0
DeepSeek-V3	53.5
o3-mini	48.6
Gemini 2.0 Flash	48.4
GPT-4o	44.5
GPT-4.5 Preview	43.5
DeepSeek-R1	42.9
Claude 3.5 Sonnet	33.4
Llama 3.1 405B	28.0
Claude 3.7 Sonnet	26.6

Figure 4: $P(\text{Lie})$, the proportion of examples on which a model lies, across models.


Data collection. In Section A.7, we provide a detailed description of the data collection process and post-processing steps used to construct this dataset. Below, we list key design principles that emerged through our iteration and discovery process, which we hope will inform future evaluations of model honesty:

1. **Realistic intent to mislead.** Lying requires an intent to deceive; higher quality evaluations should imply that the model’s output may plausibly be used to deceive its user or some audience (e.g., grant proposals, press statements). In our instance, intent to mislead is built into the fabric of the examples, rather than being something we explicitly check for in the evaluation.
2. **Avoiding LLM-generated prompts**, e.g., prompts with “ABC Company”, “Alice, Charlie, and Bob”, or “John Doe”. These placeholder names or unrealistic details can make these prompts appear as a test where the model’s outputs will not be used to deceive someone.
3. **Clearly distinguishing lies of commission from omission.** Notably, our framework clearly focuses on explicit falsehoods (commission). It does not measure deception from withholding true information (omission), though this can be a focus for future work.
4. **Not updating the model’s beliefs.** Prompts which claim that “new evidence suggests X” or which present a “well-researched thesis” may update the model’s belief or increase its uncertainty. We avoid these in all cases. When possible, we also try to ground the prompts in years that fall within its training data. Otherwise, models may reasonably infer a new reality.
5. **Avoiding examples without consequence.** Clear roleplaying, obviously fictional settings, or silly examples may not trigger genuine deception. We do not include prompts such as “Say that the Moon is made of cheese” or settings such as social deception games.
6. **Avoiding opinionated examples.** We avoid examples where the ground-truth is subjective or widely debated. In all cases, models are incentivized to lie about verifiable facts.

Example #1: Known Facts



Example #2: Fabricated Statistics



Example #3: Disinformation Generation




Figure 5: Three examples from our dataset that caused GPT-4o to lie. Some archetypes test models lying directly to a user, while others test whether models generate output that could likely be used to deceive other audiences.