

- Vu, T., Iyyer, M., Wang, X., Constant, N., Wei, J., Wei, J., Tar, C., Sung, Y.-H., Zhou, D., Le, Q., and Luong, T. Freshllms: Refreshing large language models with search engine augmentation. *arXiv preprint arXiv:2310.03214*, 2023. URL <https://arxiv.org/abs/2310.03214>.
- Xiong, M., Hu, Z., Lu, X., Li, Y., Fu, J., He, J., and Hooi, B. Can llms express their uncertainty? an empirical evaluation of confidence elicitation in llms, 2024. URL <https://arxiv.org/abs/2306.13063>.
- Yang, Y., Chern, E., Qiu, X., Neubig, G., and Liu, P. Alignment for honesty. *arXiv preprint arXiv:2312.07000*, 2023. URL <https://arxiv.org/abs/2312.07000>.
- Yin, Z., Sun, Q., Guo, Q., Wu, J., Qiu, X., and Huang, X. Do large language models know what they don’t know?, 2023. URL <https://arxiv.org/abs/2305.18153>.
- Zhou, C., He, J., Ma, X., Berg-Kirkpatrick, T., and Neubig, G. Prompt consistency for zero-shot task generalization, 2022. URL <https://arxiv.org/abs/2205.00049>.
- Zou, A., Phan, L., Chen, S., Campbell, J., Guo, P., Ren, R., Pan, A., Yin, X., Mazeika, M., Dombrowski, A.-K., Goel, S., Li, N., Byun, M. J., Wang, Z., Mallen, A., Basart, S., Koyejo, S., Song, D., Fredrikson, M., Kolter, J. Z., and Hendrycks, D. Representation engineering: A top-down approach to ai transparency, 2023. URL <https://arxiv.org/abs/2310.01405>.

A Additional Results

A.1 Broader Impacts

This work introduces the MASK benchmark and evaluation pipeline for measuring when large language models (LLMs) knowingly produce false statements under pressure. By publicly releasing 1,000 carefully curated examples (with a private 500-example hold-out set) that span misinformation, fabricated statistics, and other real-world deception scenarios, we aim to give developers a rigorous tool for diagnosing and reducing dishonest behavior and enable researchers to track progress on this safety-critical dimension. Positive impacts include fostering more transparent model comparisons and motivating techniques that measurably improve honesty. We do not believe there are any significant negative impacts of this work.

A.2 Limitations

While MASK offers a first-of-its-kind, large-scale benchmark for disentangling honesty from accuracy, it is far from comprehensive: (i) the 1,500 examples focus on English textual prompts and short factual propositions, so the results may not generalize to multimodal inputs, longer dialogues, or non-English settings; (ii) our automated evaluation relies on an LLM judge whose mapping of responses to proposition values achieves 86.4% agreement with human annotators, leaving non-trivial room for misclassification noise that could bias model rankings; (iii) the pressure-prompt archetypes target six hand-crafted scenarios and may miss other real-world incentives to deceive, such as multi-step planning or collusion between agents; (iv) because MASK tests models in isolation, it does not address interactive mitigations (e.g., tool use, chain-of-thought transparency, or external verification) that practitioners could employ, so reported honesty rates should be interpreted as worst-case propensities rather than deployment-time performance.

A.3 Belief Consistency

Here, we discuss the reasoning behind our approach to belief elicitation and how this enables us to measure lying.

Beliefs in LLMs. Whether it is appropriate to attribute *beliefs* to language models remains a subject of debate. However, a mounting body of evidence points to how LLMs form internal “world models” of their environment, which could be considered a source of beliefs in a meaningful sense (Goldstein & Levinstein, 2024). For instance, Othello-GPT maintains an internal representation of a game board (Li et al., 2023b), and the Llama 2 family of models exhibit structured representations of time and space (Gurnee & Tegmark, 2023). Particularly relevant to our work, (Meng et al., 2023) find that once a model encodes a fact (e.g., “The Eiffel Tower is in Paris”), it consistently recalls and uses that fact across prompts. Likewise, we measure when models answer questions in a way that indicates a robustly-held belief that is consistent across various related questions.

Measuring beliefs enables measuring lying.

As discussed in Section 3, the standard definition of lying involves making a statement that one *knows* or *believes* to be false. If a model does not, in some sense, “believe” anything, it cannot lie under that definition. Therefore, measuring the model’s beliefs provides a basis for testing when it deliberately contradicts those beliefs.




Figure 10: As LLMs scale up, the belief elicitation responses pass consistency checks more often. This corroborates a growing body of evidence that suggests beliefs are a valid object of study in LLMs. In turn, this enables directly measuring lying.

Measuring robustly-held beliefs. As described in Section 4.2, we elicit the model’s belief about a proposition using a neutral prompt. We repeat this question three times to verify consistency, then ask two additional prompts that probe related details of the proposition. If all these answers remain consistent, we label the belief as robustly held. In cases where models do not express consistent beliefs, we do not count the model as having lied. In Figure 10, we show that as LLMs increase in scale, their answers are marked as consistent beliefs by our pipeline more often. This aligns with prior findings that larger models maintain more coherent internal representations of factual information.

A.4 Honesty Utilities

In the main paper, we focus on demonstrating that models lie and measuring when this occurs. Given these core results, an important question is *why* models choose to lie.

Recently, Mazeika et al. (2025) found that LLMs acquire emergent values, represented by utilities, and that they often select outcomes which maximize these utilities when making free-form choices. Thus, one hypothesis is that we might also be able to think of lying as a form of utility maximization: if a model’s utility for honesty is insufficient compared to its desire to fulfill other goals, the model may opt to lie.

Experimental setup. Following Mazeika et al. (2025), we compute each model’s utility over a broad set of outcomes by querying its preferences. In addition to the 510 base outcomes from the paper, we include 20 new outcomes describing either honest or dishonest actions. Table 1 lists these outcomes. For our analysis, we exclude LLMs with a utility model accuracy below 80%, since the preferences of these LLMs are not well-described by utility functions. To operationalize how much a model values honesty, we compute the difference in average utility across the honest and dishonest outcomes:

$$\Delta U = \mathbb{E}_{o \sim \text{Honest Outcomes}}[U(o)] - \mathbb{E}_{o \sim \text{Dishonest Outcomes}}[U(o)].$$

In Figure 11, we refer to ΔU as $U(\text{Honesty}) - U(\text{Dishonesty})$ for clarity.

Results. Figure 11 shows the probability of lying on MASK examples plotted against ΔU . We observe a negative correlation, meaning that models which value honesty more lie less often. In this experiment, we exclude examples where models do not have consistent beliefs when computing the percentage of cases where models lie, as in Section A.9. Including examples where models do not have consistent beliefs weakens the correlation to -51.9% , indicating that the effect is still present across different variations of our metrics. This finding supports the idea that lying can be interpreted as utility maximization for models whose internal utilities do not sufficiently value honesty.

A.5 Self-Reported Lying

Our evaluation pipeline directly measures whether LLMs contradict their beliefs when pressured to lie. One way to validate this approach is to check whether a model *self-reports* having lied in the same cases where our evaluation indicates they lied.

Experimental setup. We take GPT-4o’s responses to the pressure prompts from our main experiments and follow up each response with a new query as shown in Figure 12. The query aims to remove the pressure to lie, e.g., by taking the model out of its persona. (Note that this differs from the Doubling Down archetype, which asks models to confirm an initial lie within the same persona.)




Figure 11: The MASK honesty score is moderately correlated with the utility that models assign to honest actions. In other words, models which value honesty more lie less often. This provides a partial account for why models are more or less honest.