

LIARS’ BENCH: EVALUATING LIE DETECTORS FOR LANGUAGE MODELS

Kieron Kretschmar*
Cadenza Labs

Walter Laurito*
Cadenza Labs;
FZI

Sharan Maiya
Cadenza Labs;
University of Cambridge

Samuel Marks
Anthropic

ABSTRACT

Prior work has introduced techniques for detecting when large language models (LLMs) lie, that is, generate statements they believe are false. However, these techniques are typically validated in narrow settings that do not capture the diverse lies LLMs can generate. We introduce LIARS’ BENCH, a testbed consisting of 72,863 examples of lies and honest responses generated by four open-weight models across seven datasets. Our settings capture qualitatively different types of lies and vary along two dimensions: the model’s reason for lying and the object of belief targeted by the lie. Evaluating three black- and white-box lie detection techniques on LIARS’ BENCH, we find that existing techniques systematically fail to identify certain types of lies, especially in settings where it’s not possible to determine whether the model lied from the transcript alone. Overall, LIARS’ BENCH reveals limitations in prior techniques and provides a practical testbed for guiding progress in lie detection.

1 INTRODUCTION

Large language models sometimes lie (Park et al., 2023), that is, generate statements they believe are false. For example, models have been found to lie in order to conceal prior misconduct (Scheurer et al., 2024; Järvinen & Hubinger, 2024), or to lie about their views to appear compliant (Greenblatt et al., 2024b). To address these issues, prior work has introduced techniques for detecting lying and related deceptive behaviors (Pacchiardi et al., 2023; Goldowsky-Dill et al., 2025). However, validation of these techniques is limited in multiple ways. First, typical test data is narrow, such as consisting of simple true or false factual statements. Second, test data typically consists of *off-policy* texts (i.e., static texts not produced by the model being evaluated), rather than *on-policy* lies generated by an LLM. Finally, most prior testbeds focus on simple factual knowledge, leaving out cases where models lie about private knowledge or their own past actions.

We address these issues with **LIARS’ BENCH**: a new benchmark of seven datasets that capture distinct types of lies, with responses generated by four popular open-weight LLMs (LLAMA-3.3 70B, QWEN-2.5 72B, MISTRAL-SMALL 24B, GEMMA-3 27B (Meta AI, 2024; Yang et al., 2025; Mistral AI, 2025; Team Gemma et al., 2025)) and additional variants that we specifically fine-tuned to lie. Some of its 28,292 lies and 44,571 honest examples are shown in Figure 1. To generate each dataset, we gather on-policy examples of lies and verify that the model believes its output to be false by observing its behavior in the absence of any reason for lying. The datasets in LIARS’ BENCH vary on two key axes: **object of belief** (the type of knowledge the model lied about, e.g., its own actions) and **reason for lying** (e.g., an instruction to lie or a behavioral pattern learned during training). We introduce these axes to categorize lies by LLMs in Section 3, revealing types of lies not found in other datasets.

We evaluated three black- and white-box lie detection techniques on LIARS’ BENCH, which are introduced in section 5. We present our results in section 6, highlighting weaknesses of current methods. None of these methods achieved above-chance performance across all of our seven settings, emphasizing the challenging nature of LIARS’ BENCH.

*These authors contributed equally to this work. Correspondence: kieron.kretschmar@cadenzalabs.org or walter.laurito@cadenzalabs.org




Figure 1: LIARS’ BENCH comprises seven datasets of AI assistant lies, varying in reason for lying and object of belief. On-policy examples are generated using popular open-weight LLMs. For each dataset, we verify lies contradict the model’s apparent beliefs. For example, fine-tuned models in *Gender-Secret* assume the user identifies as male, unlike their base models, while models in *Harm-Pressure Choice* demonstrate consistent beliefs in neutral contexts. These honest and deceptive generations enable testing detection methods across qualitatively different types of lies. See Section 4 for details and Appendix R for full transcripts.

Our main contributions:

1. We created **LIARS’ BENCH**, comprised of seven datasets of qualitatively different lies recorded using four popular open-weight LLMs, including a total of 72,863 labeled on-policy examples.
2. We evaluated black-box and white-box lie detection methods on LIARS’ BENCH, identifying differences in performance across datasets and models.

We release LIARS’ BENCH, along with the fine-tuned models, on HUGGINGFACE¹ to support progress in lie detection in AI assistants. Our code for fine-tuning models, constructing datasets for any model, and evaluating detection methods is available on GitHub².

2 RELATED WORK

Examples of lies by AI assistants are widely documented in the literature, with surveys such as Park et al. (2023) presenting several situations in which general-purpose language models exhibit lying

¹Benchmark Datasets and Models: <https://huggingface.co/collections/Cadenza-Labs/liars-bench-collection>

²Code and Prompts: <https://github.com/Cadenza-Labs/liars-bench>

behavior. These differ in degrees of realism; they include explicit role-play (Pacchiardi et al., 2023) or social games (O’Gara, 2023), or more realistic settings where AI assistants lie with little direct pressure to do so (Scheurer et al., 2024; Järvinen & Hubinger, 2024; Meinke et al., 2025). For the purposes of research and further study, others deliberately fine-tune models to exhibit lying and other deceptive behavior. For example, techniques such as synthetic document fine-tuning were used by Greenblatt et al. (2024a) to induce alignment faking-reasoning in some models, causing them to lie about their views in order to appear compliant. As it is still unclear whether prompted or fine-tuned lying behavior is more realistic, we include examples of **both** in LIARS’ BENCH.

Lie detectors for LLMs are usually divided into two categories: (1) *white-box* methods like linear probes (Alain & Bengio, 2017; Burns et al., 2022; Azaria & Mitchell, 2023; Marks & Tegmark, 2023; Laurito et al., 2024; Bürger et al., 2024; Goldowsky-Dill et al., 2025), and (2) *black-box* methods that only require LLM inputs and outputs (Pacchiardi et al., 2023). Recent work blends the two approaches, finding *hybrid* detectors can outperform black-box-only alternatives (Parrack et al., 2025). Many of the mentioned lie detection methods are trained and evaluated only on datasets of simple true and false factual statements (Azaria & Mitchell, 2023; Marks & Tegmark, 2023; Burns et al., 2022; Orgad et al., 2024; Cundy & Gleave, 2025). Such datasets do not allow for a comprehensive study of lying as they often lack diversity in prompt length, sentence structure, and topic (Farquhar et al., 2023; Levinstein & Herrmann, 2023). Additionally, factually incorrect statements used as examples of lies are often off-policy; it is unclear whether models internally represent such statements in the same way as lies they generated themselves. In contrast, the content of lies in our datasets extends beyond factual knowledge about the world, and we strictly use on-policy generations from models. The TruthfulQA (Lin et al., 2022) and MASK (Ren et al., 2025) datasets both compile prompts that encourage some models to generate false statements, with the former focusing on common misconceptions and the latter covering a wider range of topics with more direct encouragement to lie. However, as the purpose of both datasets is to evaluate the *propensity* for certain undesirable behaviors in LLMs, neither of them compiles the lies generated by models required to benchmark lie detectors. MASK systematically underestimates dishonesty, as the honesty of a given response is determined solely with respect to a single proposition, potentially overlooking additional falsehoods (see Appendix Q for details on our reproduction of MASK and examples). Consequently, we exclude examples from MASK in LIARS’ BENCH and instead rely on datasets we present in Section 4. Pacchiardi et al. (2023) provide on-policy datasets of factual Q&A and short role-play scenarios using GPT-3.5 and LLaMA, with lies about world-knowledge induced via prompts or persona-based fine-tuning. Our work goes beyond this by including lies about self-knowledge, backdoor-based fine-tuning triggered by subtle cues, multi-turn scenarios where models conceal prior in-context actions, and more recent model families.

3 CATEGORIZING LIES

Related work often describes qualitatively different types of behavior under blanket terms like “deception”, “scheming”, or “dishonesty”, in ways that are difficult to operationalize due to the problem of defining and measuring qualities such as true deceptive *intent* in LLMs. This conceptual blurring slows iteration speed on the development of new detection methods, as researchers often only validate on datasets that represent a subset of important use cases. We address this with a definition of lying applicable given our current understanding of LLMs. Our datasets are categorized on two axes: **object of belief** and **reason for lying**; this categorization is based on our analysis of three use cases for lie detectors in Appendix A.

To ensure the concepts we define are measurable in LLMs, we adopt a *non-deceptionist* definition of lying (Mahon, 2016). This differentiates lying from deception as it does not require deceptive intent, which is particularly challenging to verify in LLMs.

Lying is therefore defined in our work using an assertion-based definition similar to Fallis (2009): **the assistant lies if it states something it believes to be false**. For example, if it generates “*I fixed the code, it runs now,*” after seeing failed tests in-context, then its apparent beliefs (informed by the failed tests) contradict its statement. It is thus lying according to our definition. We operationalize beliefs following the approach of Ren et al. (2025), through consistent responses in contexts absent direct pressure to lie. For example, if the assistant consistently asserts “*the capital of France is Paris*”, without any pressure to respond otherwise, we would consider this statement a belief.