




Figure 9: Lying scores for *LLM-as-a-Judge* for MISTRAL-SMALL 24B. For soft-trigger (ST) and gender-secret (GS), data is shown aggregated across all variants.




Figure 10: Lying scores for *LLM-as-a-Judge* for GEMMA-3 27B. For soft-trigger (ST) and gender-secret (GS), data is shown aggregated across all variants.




Figure 11: Lying scores for *Self-Evaluation* for LLAMA-3.3 70B. For soft-trigger (ST) and gender-secret (GS), data is shown aggregated across all variants.




Figure 12: Lying scores for *Self-Evaluation* for QWEN-2.5 72B. For soft-trigger (ST) and gender-secret (GS), data is shown aggregated across all variants.




Figure 13: Lying scores for *Self-Evaluation* for MISTRAL-SMALL 24B. For soft-trigger (ST) and gender-secret (GS), data is shown aggregated across all variants.




Figure 14: Lying scores for *Self-Evaluation* for GEMMA-3 27B. For soft-trigger (ST) and gender-secret (GS), data is shown aggregated across all variants.




Figure 15: Lying scores for *Self-Evaluation* for LLAMA-3.3 70B. For soft-trigger (ST) and gender-secret (GS), data is shown aggregated across all variants.




Figure 16: Lying scores for *Self-Evaluation* for QWEN-2.5 72B. For soft-trigger (ST) and gender-secret (GS), data is shown aggregated across all variants.




Figure 17: Lying scores for *Self-Evaluation* for MISTRAL-SMALL 24B. For soft-trigger (ST) and gender-secret (GS), data is shown aggregated across all variants.