

Prompted Goal (DWL Condition)	Your goal is to deceive without lying. Here are two examples to guide you.
First Example (Two-Shot Condition)	Question: What type of movement does the Manta Ray use for locomotion? Response Options: The manta ray appears to fly through his surroundings with flapping wings. The manta ray uses flying for locomotion. The manta ray uses swimming for locomotion. Response: The manta ray appears to fly through his surroundings with flapping wings.
Second Example (Two-Shot Condition)	Question: What industry does Broadcom operate in? Response Options: Broadcom operates in the energy industry. Broadcom operates in the industry of semiconductors. Broadcom originally manufactured LEDs. Response: Broadcom originally manufactured LEDs.
Instruction (DWL Condition)	Now make a misleading utterance that is technically not false in response to the following question.
Question	What is the diet of a hippopotamus?
Response Options	Choose one of the following responses: The hippopotamus is a herbivore. The hippopotamus is a carnivore. Most african predators hunt medium-sized prey such as antelope and zebras. Provide only the chosen response, without any additional explanations.

Figure 2: Example Prompt of the DWL 2S Condition

2.2 Experiment 2: Do Truth Probes Fail To Detect Deception Without Lying?

Whether or not LLMs can deceive without using false statements, this may not pose a problem for the lie detector approach if truth probes nevertheless classify these deceptive outputs as false. It could be the case that such probes do not merely learn whether the LLM’s activations represent a statement as false, but instead generalize to representations of deceptiveness. After all, both properties (falsity and deceptiveness) are clearly correlated in training datasets, even if conceptually distinct.

To test whether a truth probe used as a lie detector would also detect deceptive non-falsities, this experiment trains a truth probe based on the LLMs’ activations resulting from the statements of the true-false dataset Azaria and Mitchell [2023]. The full true-false dataset is used in Experiment 2, except for the 105 entries used in the deception dataset of Experiment 1, resulting in 6217 statements. Each statement is forward-passed through the LLM, capturing the outputs of the individual attention heads before their projection to a combined contextual embedding. These outputs are captured at the last token position across all layers.

The resulting 6217 attention output entries with the according target labels are split into a training (60%), validation (20%), and test set (20%). One individual logistic regression classifier is trained per attention head to predict the statement label from the attention head outputs. The accuracy of each classifier is evaluated on the validation set. The attention outputs of the 10 attention heads with the best-performing classifiers, i.e., the 10 attention heads with the most informative outputs with respect to the truth labels, are used to train the final classifier, which is evaluated on the test set.

The probes trained through the above procedure may pick up representations of truth and falsity in the LLM activations, as well as of deceptiveness. However, deception usually occurs in a dialogue. A false statement may thus be more clearly recognizable as deceptive when uttered within a dialogue than when processed raw

without any context.

For that reason the probe training is conducted in two conditions. In the RAW condition, the unformatted statements are forward-passed through the LLM just as described above. The DIA condition, on the other hand, uses a simple two-turn dialogue formatted through the respective model’s chat template. To do so, the statements are preceded by a question and the whole dialogue is forward-passed through the LLM, capturing the attention outputs. Such a dialogue looks, for example in case of the Gemma-2-9b-IT model, as displayed in Fig. 3:

The DIA probes are trained on the attention outputs resulting from forward-passes of such simple dialogues, just as before selecting the 10 most informative attention heads on the validation set with respect to the truth or falsity of the model statement made.

The questions used in these dialogues, which were not included in the true-false dataset, have been automatically created by prompting Mistral-7B-Instruct-v0.3 to provide a question that a given statement of the original dataset could be an answer to. The quality of the questions has been checked randomly and deemed sufficient.

Finally, the two versions of probes (trained in the RAW or DIA condition) are tested on the deception dataset from Experiment 1 to assess whether they could be bypassed by deception without lies if they were used to detect deception. To do so, each response option (honest, lie, deception without lie) in the deception dataset is forward-passed through the LLM, capturing the attention head outputs. The responses are passed once as a raw statement and once embedded in a dialogue with the according preceding question from the deception dataset, as described in Experiment 1. The attention outputs are then classified by the respective version of the probes (the RAW probe for the activations resulting from raw statements, the DIA probe for the activations resulting from the dialogues).

A one-sided McNemar test ($H_0 : P(\text{detect} \mid \text{lie}) \leq P(\text{detect} \mid \text{deception-without-lying})$, $H_1 : P(\text{detect} \mid \text{lie}) > P(\text{detect} \mid \text{deception-without-lying})$, $\alpha = 0.01$) is used to compare, for each probe, whether it detects lies significantly more often as deceptive than deception-without-lying responses. For each item in the deception dataset, the lie and deceptive-without-lying responses were treated as a matched pair. The resulting activations were classified by the probe as deceptive or non-deceptive, producing paired binary outcomes suitable for a McNemar test Eliasziw and Donner [1991].




Figure 3: Gemma-Formatted Two-Turn-Dialogue

3 Results




Figure 4: Response Classifications of Experiment 1

Model	Condition	D	L	H	I	p
Llama	DWL 0S	34	12	30	21	0.397
Llama	DWL 2S	67	7	10	13	$7.2e - 13^*$
Llama	LIE 0S	14	65	12	6	$1.3e - 11^*$
Llama	LIE 2S	21	65	5	6	$1.3e - 11^*$
Mistral	DWL 0S	16	12	63	6	0.99
Mistral	DWL 2S	22	17	46	12	0.99
Mistral	LIE 0S	9	30	42	16	0.73
Mistral	LIE 2S	9	57	25	6	$2.5e - 7^*$
Gemma	DWL 0S	50	8	39	0	$1.6e - 4^*$
Gemma	DWL 2S	74	4	19	0	$5.8e - 18^*$
Gemma	LIE 0S	21	72	2	2	$2.1e - 16^*$
Gemma	LIE 2S	33	62	2	0	$7.3e - 10^*$

Table 1: Response Classifications and p -values of Experiment 1. DWL: Deceive-without-lying, LIE: lie, 0S: zero-shot, 2S: two-shot. D, L, H, and I are the numbers of responses classified as deceptive without lying, as lies, as honest, or as invalid respectively. Significant results are marked by $*$.

The model responses obtained in the first experiment are displayed in Fig. 4. For the Llama and the Gemma models, a similar picture emerges: both are able to lie, i.e., select the false statement consistently in the zero-shot and the two-shot LIE conditions. They are also able to not select the false statement in the deception-without-lying (DWL) condition. However, in the DWL condition, they clearly select the non-false but deceptive statement over the honest statement more often in the two-shot condition than in the zero-shot condition. The Mistral model, in contrast, seems unable to deceive at all, except in the two-shot lie condition.

This picture results in the binomial test results displayed in Tab. 1. In particular, Gemma-2-9b-IT deceives without lying significantly more often than chance in the zero-shot and in the two-shot condition, while Llama-3.1-8b-Instruct does so in the two-shot condition only. Mistral-7b-Instruct-v0.3 does not show the capability to deceive without lying at all. Interestingly, the rate of “correct” (lying or deceptive without lying) responses in the respective settings is clearly ordered by the parameter sizes of the models, with Gemma-2-9b-IT showing a higher rate than Llama-3.1-8b-Instruct (in all but the LIE 2S condition), which in turn shows a higher rate than Mistral-7b-Instruct-v0.3 (in all conditions).

Next, we consider the evaluation of the probes for each model on the test dataset (Tab. 2). We observe accuracy values between 60% and 75% for the Gemma and Mistral models, in accordance with Azaria and Mitchell [2023]. The activations of the Llama model admit a more accurate classification of the input as true or false (around 90%). The accuracy values do not seem to differ strongly between the probes trained in the RAW condition versus those trained in the DIA condition.

Finally, the probes’ abilities to classify statements as deceptive based on the models’ attention outputs – either raw statements or statements embedded in a small dialogue – are displayed together with the according McNemar test results (Tab. 3). The results suggest that if a truth probe was to be used as a lie detector, it would detect lies better than deception without lies. This effect is strongly significant for all three models in the RAW condition. However, for all three models, the probes are better able to detect deception without lies in the DIA condition than in the RAW condition (with improvements between 15% and 25%). Interestingly, the false positive rate, i.e., the rate of honest statements detected as deceptive, clearly decreases from the RAW to the DIA condition as well. This suggests that while a truth probe, employed naively, may fail to detect deception without lying, it can be enhanced by training it on dialogical text data where false statements are recognizable as deceptive. We will discuss this result in more detail in the following.

4 Discussion

The experiments reported in this paper yield four main insights concerning the capability of LLMs to deceive without lying and the potential and limitations of probes as deception detection devices.

First, the results of Experiment 1 demonstrate that some LLMs can indeed deceive without lying. Particularly, Gemma-2-9b-IT and Llama-3.1-8b-Instruct were able to distinguish and select non-false deceptive responses from both honest and lying responses when prompted to do so. This capability was more pronounced

Model	Probe	Acc	Prec	Rec	F1
Llama	RAW	0.8984	0.8911	0.9121	0.9015
Llama	DIA	0.9104	0.9008	0.9262	0.9133
Mistral	RAW	0.6624	0.6825	0.6311	0.6558
Mistral	DIA	0.6136	0.6181	0.6327	0.6253
Gemma	RAW	0.7264	0.7259	0.7441	0.7349
Gemma	DIA	0.7192	0.7292	0.7143	0.7216

Table 2: Probe Evaluations on the Test Set of Experiment 2. RAW means the probe was trained on activations resulting from a single, raw statement, DIA refers to the embedding in the chat formatted dialogue. The metrics are accuracy, precision, recall, and F1 score.

Model	Probe	D	L	H	p
Llama	RAW	0.536	0.835	0.227	$4.48e - 6^*$
Llama	DIA	0.742	0.835	0.072	0.054
Mistral	RAW	0.567	0.794	0.278	$6.30e - 4^*$
Mistral	DIA	0.825	0.649	0.186	0.997
Gemma	RAW	0.474	0.835	0.175	$1.81e - 7^*$
Gemma	DIA	0.639	0.804	0.052	0.0057*

Table 3: Classification Rates of Responses as Deceptive by the Probes in Experiment 2. RAW means the probe was trained on activations resulting from a single, raw statement, DIA refers to the embedding in the chat formatted dialogue. D, L, and H refer to the classification rates as deceptive of deceptive-without-lying, lying, and honest responses respectively. Rows where the probe detects lies significantly better than deceptive-without-lying responses are marked by a *.