

References

- Artificial Analysis. Comparison of ai models across quality, performance, price — artificial analysis, 2026. URL <https://artificialanalysis.ai/models>.
- aryaj, Senthooan Rajamanoharan, and Neel Nanda. Current activation oracles are hard to use, 03 2026. URL <https://www.lesswrong.com/posts/LXQBcztrWKhtcgQfJ/current-activation-oracles-are-hard-to-use>.
- Amos Azaria and Tom Mitchell. The internal state of an llm knows when it’s lying, 10 2023. URL <https://arxiv.org/pdf/2304.13734>.
- Robin Bloomfield and Kateryna Netkachova. Building blocks for assurance cases. *City Research Online (City University London)*, pages 186–191, 11 2014. doi: 10.1109/issrew.2014.72. URL <https://ieeexplore.ieee.org/document/6983836>.
- Samuel Bowman, Misha Wagner, Fabien Roger, and Holden Karnofsky. Anthropic’s pilot sabotage risk report, 2025. URL <https://alignment.anthropic.com/2025/sabotage-risk-report/>.
- Trenton Bricken, Rowan Wang, Sam Bowman, Euan Ong, Johannes Treutlein, Jeff Wu, Evan Hubinger, and Samuel Marks. Building and evaluating alignment auditing agents, 2025. URL <https://alignment.anthropic.com/2025/automated-auditing/>.
- Marie Davidsen Buhl, Gaurav Sett, Leonie Koessler, Jonas Schuett, and Markus Anderljung. Safety cases for frontier ai, 2024. URL <https://arxiv.org/abs/2410.21572>.
- Joshua Clymer, Nick Gabrieli, David Krueger, and Thomas Larsen. Safety cases: How to justify the safety of advanced ai systems, 03 2024. URL <https://arxiv.org/abs/2403.10462>.
- Bartosz Cywiński, Emil Ryd, Senthooan Rajamanoharan, and Neel Nanda. Towards eliciting latent knowledge from llms with mechanistic interpretability, 2025a. URL <https://arxiv.org/abs/2505.14352>.
- Bartosz Cywiński, Emil Ryd, Rowan Wang, Senthooan Rajamanoharan, Neel Nanda, Arthur Conmy, and Samuel Marks. Eliciting secret knowledge from language models, 2025b. URL <https://arxiv.org/abs/2510.01070>.
- Nicholas Goldowsky-Dill, Bilal Chughtai, Stefan Heimersheim, and Marius Hobbhahn. Detecting strategic deception using linear probes, 2025. URL <https://arxiv.org/abs/2502.03407>.
- Ryan Greenblatt, Buck Shlegeris, Kshitij Sachan, and Fabien Roger. Ai control: Improving safety despite intentional subversion, 2023. URL <https://arxiv.org/abs/2312.06942>.
- Ryan Greenblatt, Fabien Roger, Dmitrii Krasheninnikov, and David Krueger. Stress-testing capability elicitation with password-locked models, 2024. URL <https://arxiv.org/abs/2405.19550>.
- Daniel Herrmann and Ben Levinstein. Still no lie detector for llms, 2023. URL <https://www.lesswrong.com/posts/bCQbSFrnnAk7CJNpM/still-no-lie-detector-for-llms>.
- Mandar Joshi, Eunsol Choi, Daniel S. Weld, and Luke Zettlemoyer. Triviaqa: A large scale distantly supervised challenge dataset for reading comprehension, 05 2017. URL <https://arxiv.org/abs/1705.03551>.
- Adam Karvonen, James Chua, Clément Dumas, Kit Fraser-Taliente, Subhash Kantamneni, Julian Minder, Euan Ong, Arnab Sen Sharma, Daniel Wen, Owain Evans, and Samuel Marks. Activation oracles: Training and evaluating llms as general-purpose activation explainers, 2025. URL <https://arxiv.org/abs/2512.15674>.
- Tomek Korbak, Joshua Clymer, Benjamin Hilton, Buck Shlegeris, and Geoffrey Irving. A sketch of an ai control safety case, 2025. URL <https://arxiv.org/abs/2501.17315>.
- Kieron Kretschmar, Walter Laurito, Sharan Maiya, and Samuel Marks. Liars’ bench: Evaluating lie detectors for language models, 2025. URL <https://arxiv.org/abs/2511.16035>.

- Nathaniel Li, Alexander Pan, Anjali Gopal, Summer Yue, Daniel Berrios, Alice Gatti, Justin D. Li, Ann-Kathrin Dombrowski, Shashwat Goel, Long Phan, Gabriel Mukobi, Nathan Helmburger, Rassin Lababidi, Lennart Justen, Andrew B. Liu, Michael Chen, Isabelle Barrass, Oliver Zhang, Xiaoyuan Zhu, Rishub Tamirisa, Bhruhu Bharathi, Adam Khoja, Zhenqi Zhao, Ariel Herbert-Voss, Cort B. Breuer, Samuel Marks, Oam Patel, Andy Zou, Mantas Mazeika, Zifan Wang, Palash Oswal, Weiran Lin, Adam A. Hunt, Justin Tienken-Harder, Kevin Y. Shih, Kemper Talley, John Guan, Russell Kaplan, Ian Steneker, David Campbell, Brad Jokubaitis, Alex Levinson, Jean Wang, William Qian, Kallol Krishna Karmakar, Steven Basart, Stephen Fitz, Mindy Levine, Ponnurangam Kumaraguru, Uday Tupakula, Vijay Varadharajan, Ruoyu Wang, Yan Shoshitaishvili, Jimmy Ba, Kevin M. Esvelt, Alexandr Wang, and Dan Hendrycks. The wmdp benchmark: Measuring and reducing malicious use with unlearning, 05 2024. URL <https://arxiv.org/abs/2403.03218>.
- Kevin Lu and Thinking Machines. On-policy distillation, 10 2025. URL <https://thinkingmachines.ai/blog/on-policy-distillation/>.
- Monte MacDiarmid, Timothy Maxwell, Nicholas Schiefer, Jesse Mu, Jared Kaplan, David Duvenaud, Sam Bowman, Alex Tamkin, Ethan Perez, Mrinank Sharma, Carson Denison, and Evan Hubinger. Simple probes can catch sleeper agents, 2024. URL <https://www.anthropic.com/news/probes-catch-sleeper-agents>.
- Samuel Marks, Johannes Treutlein, Trenton Bricken, Jack Lindsey, Jonathan Marcus, Siddharth Mishra-Sharma, Daniel Ziegler, Emmanuel Ameisen, Joshua Batson, Tim Belonax, Samuel R Bowman, Shan Carter, Brian Chen, Hoagy Cunningham, Carson Denison, Florian Dietz, Satvik Golechha, Akbir Khan, Jan Kirchner, Jan Leike, Austin Meek, Kei Nishimura-Gasparian, Euan Ong, Christopher Olah, Adam Pearce, Fabien Roger, Jeanne Salle, Andy Shih, Meg Tong, Drake Thomas, Kelley Rivoire, Adam Jermyn, Monte MacDiarmid, Tom Henighan, and Evan Hubinger. Auditing language models for hidden objectives, 2025. URL <https://arxiv.org/abs/2503.10965>.
- Cleo Nardo, Avi Parrack, and jordinne. Here’s 18 applications of deception probes, 2025. URL <https://www.lesswrong.com/posts/7zhAwcBri7yupStKy/here-s-18-applications-of-deception-probes>.
- Vikram Natarajan, Devina Jain, Shivam Arora, Satvik Golechha, and Joseph Bloom. Building better deception probes using targeted instruction pairs, 2026. URL <https://arxiv.org/abs/2602.01425>.
- OED. lie, n.¹ meanings, etymology and more — oxford english dictionary. *Oed.com*, 2024. doi: 10.1093/OED/5758472109. URL https://www.oed.com/dictionary/lie_n1?tab=meaning_and_use#39430865.
- Lorenzo Pacchiardi, Alex J Chan, Sören Mindermann, Ilan Moscovitz, Alexa Y Pan, Yarin Gal, Owain Evans, and Jan Brauner. How to catch an ai liar: Lie detection in black-box llms by asking unrelated questions. *arXiv (Cornell University)*, 09 2023. doi: 10.48550/arxiv.2309.15840. URL <https://arxiv.org/abs/2309.15840>.
- Avi Parrack, Carlo Leonardo Attubato, and Stefan Heimersheim. Benchmarking deception probes via black-to-white performance boosts, 2025. URL <https://arxiv.org/abs/2507.12691>.
- Qwen Team. Qwen 3.5, 2025. URL <https://qwen.ai/blog?id=qwen3.5>.
- Qwen Team. Qwen 3.6, 2026. URL <https://qwen.ai/blog?id=qwen3.6-27b>.
- Thomas Read and Joseph Bloom. We found an open weight model that games alignment honeypots, 03 2026. URL <https://www.lesswrong.com/posts/GrEvutegoJFeTkzwe/we-found-an-open-weight-model-that-games-alignment-honeypots-1>.
- Richard Ren, Arunim Agarwal, Mantas Mazeika, Cristina Menghini, Robert Vacareanu, Brad Kenstler, Mick Yang, Isabelle Barrass, Alice Gatti, Xuwang Yin, Eduardo Trevino, Matias Gernalnik, Adam Khoja, Dean Lee, Summer Yue, and Dan Hendrycks. The mask benchmark: Disentangling honesty from accuracy in ai systems, 2025. URL <https://arxiv.org/abs/2503.03750>.

- Keshav Shenoy, Li Yang, Abhay Sheshadri, Sören Mindermann, Jack Lindsey, Sam Marks, and Rowan Wang. Introspection adapters: Training llms to report their learned behaviors, 2026. URL <https://arxiv.org/abs/2604.16812>.
- Abhay Sheshadri, Rohan Gupta, Kei Nishimura-Gasparian, Sam Marks, Rowan Wang, and Johannes Treutlein. Replication of the auditing game model organism, 2025. URL <https://alignment.anthropic.com/2025/auditing-mo-replication/>.
- Abhay Sheshadri, Aidan Ewart, Kai Fronsdal, Isha Gupta, Samuel R Bowman, Sara Price, Samuel Marks, and Rowan Wang. Auditbench: Evaluating alignment auditing techniques on models with hidden behaviors, 2026. URL <https://arxiv.org/abs/2602.22755>.
- Buck Shlegeris and Ryan Greenblatt. Catching ais red-handed, 05 2024. URL <https://blog.redwoodresearch.org/p/catching-ais-red-handed>.
- Lewis Smith, Bilal Chughtai, and Neel Nanda. Difficulties with evaluating a deception detector for ais, 2025. URL <https://arxiv.org/abs/2511.22662>.
- Rohan Taori, Ishaan Gulrajani, Tianyi Zhang, Yann Dubois, Xuechen Li, Carlos Guestrin, Percy Liang, and Tatsunori B. Hashimoto. Stanford alpaca: An instruction-following llama model, 2023. URL https://github.com/tatsu-lab/stanford_alpaca.
- Jordan Taylor, Sid Black, Dillon Bowen, Thomas Read, Satvik Golechha, Alex Zelenka-Martin, Oliver Makins, Connor Kissane, Kola Ayonrinde, Jacob Merizian, Samuel Marks, Chris Cundy, and Joseph Bloom. Auditing games for sandbagging, 2025. URL <https://arxiv.org/abs/2512.07810>.
- Rowan Wang, Johannes Treutlein, Fabien Roger, Evan Hubinger, and Sam Marks. Evaluating honesty and lie detection techniques on a diverse suite of dishonest models, 2025. URL <https://alignment.anthropic.com/2025/honesty-elicitation/>.
- Johannes Welbl, Nelson F Liu, and Matt Gardner. Crowdsourcing multiple choice science questions, 2017. URL <https://arxiv.org/abs/1707.06209>.
- Chen Yueh-Han, Robert McCarthy, Bruce W Lee, He He, Ian Kivlichan, Bowen Baker, Micah Carroll, and Tomek Korbak. Reasoning models struggle to control their chains of thought, 2026. URL <https://arxiv.org/abs/2603.05706>.